SHAHRIAR SHAYESTEH

Research Scientist Intern | Responsible AI, NLP, AI Agents | Summer 2027

State College, PA | +1 814-769-6022 | shahriar.shayesteh92@gmail.com | Website | Google Scholar | GitHub

Summary

Ph.D. candidate studying how AI systems decide what information gets through. Audits disclosure and omission when LLMs and AI agents speak for users or explain complex information.

Technical Skills

Programming: Python, SQL, Java, C++
ML and NLP: PyTorch, TensorFlow, scikit-learn, Hugging Face, pandas, NumPy
LLM Systems: LLM evaluation, tool-calling agents, LangChain, AutoGen, Model Context Protocol (MCP)
Research Methods: Experiment design, human annotation, topic modeling, fairness and privacy analysis

Research Experience

Graduate Research Assistant | Human Language Technologies Lab, Penn State | *Aug 2023 – Present*

- Published a diagnostic audit of 2,344 third-party API and action schemas, identifying interface conditions that may invite unnecessary disclosure by tool-using LLM agents.
- Designed ongoing controlled evaluations that vary tool-schema structure, guidance, and declared purpose while holding the user task and backend operation fixed.
- Built a source-grounded audit of wireless-contract explanations using 24 topics and 3,899 verified provisions; manuscript targeted to ACM FAccT 2027.
- Analyzed 59,590 matched privacy-policy pairs for a PoPETs 2027 revision; contributed to PrivaSeer's 3,967,487-policy corpus and SoAC/SoACer's 195,495-website corpus.

Research Intern | Department of Canadian Heritage | *Feb 2022 – Apr 2022*

- Applied topic modeling and sentiment analysis to qualitative data on challenges faced by Canadian artists and translated results into policy-relevant findings.

Graduate Research Assistant | NLP Laboratory, University of Ottawa | *Jan 2021 – Jun 2023*

- Developed NDAGAN, a semi-supervised text-classification model using negative data augmentation to mitigate mode collapse in GAN-BERT; evaluated it across benchmark datasets.

Selected Research Projects

How Tool Schemas Shape LLM Agent Disclosure: Structure, Guidance, and Declared Purpose — *Ongoing research.*

Easier to Read, Still Easy to Miss: Auditing Content Allocation in LLM Explanations of Wireless Contracts — *Manuscript in preparation; target venue: ACM FAccT 2027.*

Privacy, More or Less: A Large-Scale Comparison of Privacy-Policy Disclosure Patterns Across Online Sectors — *Under revision for PoPETs 2027.*

Selected Publications

- S. Shayesteh and S. Wilson. "From Conventional Web Privacy to Agentic Disclosure: How Tool Schemas May Invite LLM Oversharing." PrivateNLP @ ACL, 2026.
- S. Shayesteh et al. "SoAC and SoACer: A Sector-Based Corpus and LLM-Based Framework for Sectoral Website Classification." ACM DocEng, 2025.
- S. Wilson et al. "The PrivaSeer Project: Large-Scale Resources for Analysis of Privacy Policy Text." USENIX SOUPS, 2025.
- S. Shayesteh and D. Inkpen. "Generative Adversarial Learning with Negative Data Augmentation for Semi-Supervised Text Classification." FLAIRS, 2022.

Education

Pennsylvania State University | Ph.D. in Informatics, GPA 3.93/4.0 | Expected Fall 2027
University of Ottawa | M.Sc. in Computer Science, GPA 4.0/4.0 | Sep 2020 – Jun 2023
University of British Columbia | Non-degree CS and Statistics coursework | Sep 2017 – May 2020

Talks & Training: Guest lecture, IST 574 (scheduled Oct 2026); PrivateNLP @ ACL presenter (Jul 2026); Mila Responsible AI & Human Rights Summer School (Jun 2023)